

# HW\_04 Problem #1

## 1 Multi-Layer Perceptron [20pts]

In class we discussed the derivation of the backpropagation algorithm for neural networks. In this problem, you will train a neural network on the CIFAR10 data set. Train a Multi-Layer Perceptron (MLP) neural network on the CIFAR10 data set. This is an open ended implementation problem, but I expect that you implement the MLP with at least two different hidden layer sizes and use regularization.

- Report the classification error on the training and testing data each configuration of the neural network. For example, you should report the results in the form of a table

	Classification Error	
	training	testing
50HLN+no regularization	0.234	0.253
50HLN+ $L_2$ regularization	0.192	0.203
250HLN+no regularization	0.134	0.153
250HLN+ $L_2$ regularization	0.092	0.013

List all the parameters that you are using (i.e., number of learning rounds, regularization parameters, learning rate, etc.)

- I would suggest using Google's TensorFlow, PyTorch or Keras library to implement the MLP; however, you are free to use whatever library you'd like. If that is the case, here is a link to the data
- I recommend using a cloud platform such as Google Colab to run the code.

**This Code was written by Robert 'Quinn' Hull, and borrowed elements from several other resources:**

- Resources used to make HW\_04 Problem #1:
  - The 'load CIFAR' dataset heavily borrowed from pytorch tutorial - [https://pytorch.org/tutorials/beginner/blitz/cifar10\\_tutorial.html](https://pytorch.org/tutorials/beginner/blitz/cifar10_tutorial.html)
  - A refresher on the MLP borrowed from machine learning mastery - <https://machinelearningmastery.com/neural-networks-crash-course/>
  - Some code borrowed from medium . com about putting MLP into a pytorch - [https://medium.com/@aungkyawmyint\\_26195/multi-layer-perceptron-mnist-pytorch-463f795b897a](https://medium.com/@aungkyawmyint_26195/multi-layer-perceptron-mnist-pytorch-463f795b897a)
  - Some notes on regularization in MLP:
    - <https://cedar.buffalo.edu/~srihari/CSE574/Chap5/Chap5.5-Regularization.pdf>
  - Very helpful for avoiding a blocker with batch\_size dimensions. Be very careful with linear layers as the first parameter is batch\_size, this is different than for convolutional layers
    - <https://discuss.pytorch.org/t/valueerror-expected-input-batch-size-324-to-match-target-batch-size-4/24498>
    - <https://towardsdatascience.com/pytorch-layer-dimensions-what-sizes-should-they-be-and-why-4265a41e01fd>
  - L2 normalization is implemented in the optimizer

- <https://pytorch.org/docs/stable/optim.html>
- How to save a model object
  - [https://pytorch.org/tutorials/beginner/saving\\_loading\\_models.html](https://pytorch.org/tutorials/beginner/saving_loading_models.html)

In [1]: `%matplotlib inline`

In [2]: `import torch  
import torchvision  
import torchvision.transforms as transforms  
from torch.utils.data.sampler import SubsetRandomSampler  
import matplotlib.pyplot as plt  
import numpy as np  
import torch.nn as nn  
import torch.nn.functional as F  
import torch.optim as optim  
  
import pandas as pd`

## HW\_04\_Problem 1

### MLP

- Pseudocode:
  1. Load CIFAR Data
    - Show an image
  2. Build two MLP models
    - Multiple layer sizes
    - Use regularization
  3. For each model
    - A. Train a multi-layer perceptron (MLP) on the CIFAR10 data set
      - Report error on training data
        - w/ regularization
        - w/o regularization
      - include report of hyperparameters (epochs, regularization, learning rate)
    - B. Test a multi-layer perceptron (MLP) on the CIFAR10 data set
      - Report error on testing data
        - w/ regularization
        - w/o regularization
      - include report of hyperparameters (epochs, regularization, learning rate)

### 1. Load CIFAR Dataset

In [4]: `# # number of subprocesses to use for data loading  
# num_workers = 0  
# how many samples per batch to load  
batch_size = 12  
# # percentage of training set to use as validation  
# valid_size = 0.2`

```
# convert data to torch.FloatTensor
transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

# choose the training and testing datasets
train_data = torchvision.datasets.CIFAR10(root = './data', train = True, download=True)
test_data = torchvision.datasets.CIFAR10(root = './data', train = False, download=True)

train_loader = torch.utils.data.DataLoader(train_data, batch_size=batch_size,
                                           shuffle=True, num_workers=0)
test_loader = torch.utils.data.DataLoader(test_data, batch_size=batch_size,
                                           shuffle=False, num_workers=0)
```

Files already downloaded and verified

Files already downloaded and verified

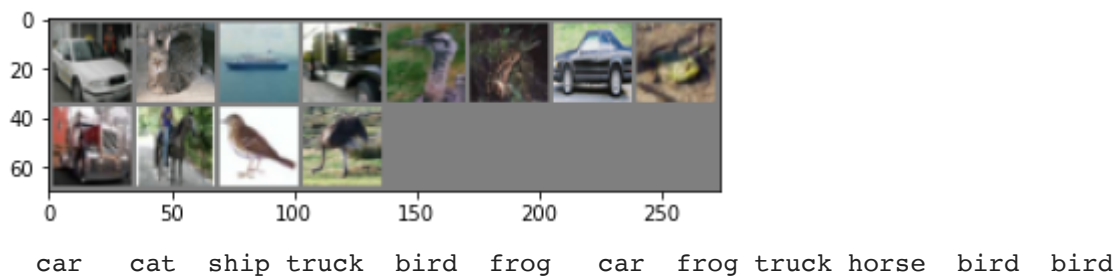
## Show an Image

```
In [5]: classes = ('plane', 'car', 'bird', 'cat',
                  'deer', 'dog', 'frog', 'horse', 'ship', 'truck')

# functions to show an image
def imshow(img):
    img = img / 2 + 0.5     # unnormalize
    npimg = img.numpy()
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
    plt.show()

# get some random training images
dataiter = iter(train_loader)
images, labels = dataiter.next()

# show images
imshow(torchvision.utils.make_grid(images))
# print labels
print(' '.join('%5s' % classes[labels[j]] for j in range(batch_size)))
```



## Data Preparation

### 2. Build an MLP model

```
In [6]: # from pytorch tutorial (see resources)
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
```

```

self.fc1 = nn.Linear(16 * 5 * 5, 120)
self.fc2 = nn.Linear(120, 84)
self.fc3 = nn.Linear(84, 10)

def forward(self, x):
    x = self.pool(F.relu(self.conv1(x)))
    x = self.pool(F.relu(self.conv2(x)))
    x = x.view(-1, 16 * 5 * 5)
    x = F.relu(self.fc1(x))
    x = F.relu(self.fc2(x))
    x = self.fc3(x)
    return x

# modified from Net2 in medium article (see resources)
class Net3(nn.Module):
    '''a two-layer linear model
        dropout: with default dropout = True
        do_const: with default dropout countant = 0.2
        hidden: with a default hidden dimension = 100
    '''

    def __init__(self, hidden=100, dropout=True, do_const=0.2):
        super(Net3, self).__init__()

        # characteristics of input
        dim_1 = 32 # x dimensions of the figure
        dim_2 = 32 # y dimensions of the figure
        num_classes = 10 # number of output classes
        num_col = 3 # since color images, for greyscale = 1

        # number of hidden nodes in each layer
        self.hidden_1 = hidden # hidden layer dimension (user defined)
        self.hidden_2 = hidden # hidden layer dimension (user defined)

        # dropout
        self.dropout = dropout
        self.do_const = do_const

        # linear layer (num_col*dim_1 * dim_2 -> hidden_1) *NOTE Adding num_col
        self.fc1 = nn.Linear(3*32*32, self.hidden_1)
        # linear layer (n_hidden -> hidden_2)
        self.fc2 = nn.Linear(self.hidden_1, self.hidden_2)
        # linear layer (n_hidden -> num_classes)
        self.fc3 = nn.Linear(self.hidden_2, 10)

        # optional dropout
        if self.dropout:
            # dropout layer (p=self.do_const)
            # dropout prevents overfitting of data
            self.dropout = nn.Dropout(p=self.do_const)
        else:
            self.dropout = nn.Dropout(p=0)

    def forward(self, x):
        # print(x.shape)
        # flatten image input (-1, num_col*dim_1*dim_2) *NOTE Adding num_col ver
        x = x.view(-1, 3*32*32)
        # print(x.shape)
        # add hidden layer, with relu activation function
        x = F.relu(self.fc1(x))
        # add dropout layer

```

```

x = self.dropout(x)
# add hidden layer, with relu activation function
x = F.relu(self.fc2(x))

# optional dropout
if self.dropout:
    # add dropout layer
    x = self.dropout(x)
else:
    x = self.dropout(x)

# add output layer
x = self.fc3(x)
return x

```

### 3. Train each model

- Prelim - Create a pandas array to store all the information re: training and testing
- Prelim - Define a loss function and optimizer

1. First, train the simplest neural net
  - with regularization
  - without regularization
2. Then, test the more complicated neural net (with convolution)
  - with regularization
  - without regularization
3. Save the data

```

In [7]: # Prelim - Create a pandas array to store all the information re: training and t
model_df = pd.DataFrame(columns=['model_name', 'net', 'train_err', 'test_error',

```

### Prelim

```

In [53]: # Prelim define loss function, optimizer, and other hyperparameters
lr = 0.001
momentum = 0.9
n_epochs = 5

# L2
L2_penalty = True # the L2 penalty is par tof the optimizer
L2_constant = 1e-2 # some L2 constant > 0

# dropout
dropout=False
do_const=0.5

# hidden layers
hidden = 50

# set up net
net_in = Net3(hidden=hidden, dropout=dropout, do_const=do_const)
print(net_in)

# controls optimization
criterion = nn.CrossEntropyLoss()
if L2_penalty: # add regularization if desired

```

```

optimizer = optim.SGD(net_in.parameters(), lr=0.001, momentum=0.9, weight_decay=0.0001)
else:
    optimizer = optim.SGD(net_in.parameters(), lr=0.001, momentum=0.9, weight_decay=0.0001)

# assembles list of loss for plotting
loss_list = []

# creates model name
model_name = str(hidden)+'HLN+'+'L2='+str(L2_penalty)
print(model_name)

Net3(
    (fc1): Linear(in_features=3072, out_features=50, bias=True)
    (fc2): Linear(in_features=50, out_features=50, bias=True)
    (fc3): Linear(in_features=50, out_features=10, bias=True)
    (dropout): Dropout(p=0, inplace=False)
)
50HLN+L2=True

```

## 4. Train the neural Net

```

In [54]: # training
for epoch in range(n_epochs): # loop over the dataset multiple times

    running_loss = 0.0 # keep track of loss within each epoch

    for i, data in enumerate(train_loader, 0): # loop through batches in train

        # get the inputs; data is a list of [inputs, labels]
        inputs, labels = data

        # zero the parameter gradients
        optimizer.zero_grad()

        # forward + backward + optimize
        outputs = net_in(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        # update loss statistics
        running_loss += loss.item()

    # print and save loss statistics
    print(running_loss / i)
    loss_list.append(running_loss / i)
    running_loss = 0.0

# final things
print('Finished Training')

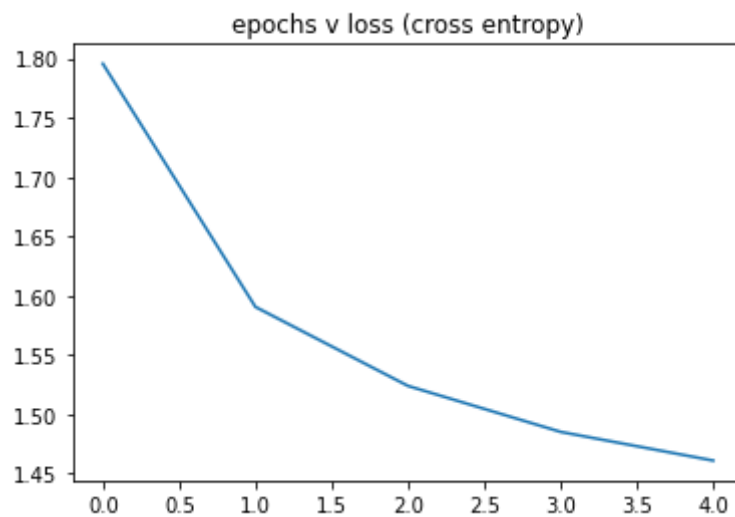
# save some results

# plot
plt.plot(np.arange(0, len(loss_list)), loss_list)
plt.title('epochs v loss (cross entropy)')
plt.show()

```

1.7955943031260673

1.5903667818309593  
 1.5236713025384492  
 1.4848844187081118  
 1.4606699741808085  
 Finished Training



## 5. test and

## 6. save

```
In [55]: # train
correct = 0
total = 0
with torch.no_grad():
    for data in train_loader:
        images, labels = data
        outputs = net_in(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

# create train error, where train error is the percent mis-identified (expressed as a percent)
train_err = 1 - (correct / total)

# test
correct = 0
total = 0
with torch.no_grad():
    for data in test_loader:
        images, labels = data
        outputs = net_in(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

# create test error, where test error is the percent mis-identified (expressed as a percent)
test_err = 1 - (correct / total)
```

```
In [58]: model_df = model_df.append({'model_name': model_name, 'net': net_in,
                                     'train_err': train_err, 'test_error': test_err, 'mod
```

PLOT

```
In [80]: model_df[['model_name', 'train_err', 'test_error', 'L2_penalty', 'num_nodes', 'num_layers']]
```

```
Out[80]:
```

	model_name	train_err	test_error	L2_penalty	num_nodes	num_layers
1	200HLN+L2=False	0.39780	0.4666	0.00	200	2
2	200HLN+L2=True	0.47864	0.5067	0.01	200	2
3	200HLN+L2=True	0.90000	0.9000	0.50	200	2
4	50HLN+L2=False	0.45254	0.4956	0.00	50	2
5	50HLN+L2=True	0.49738	0.5181	0.01	50	2
6	50HLN+L2=True	0.90000	0.9000	0.50	50	2
7	50HLN+L2=True	0.90000	0.9000	0.50	50	2
8	5HLN+L2=False	0.65572	0.6646	0.00	5	2

```
In [98]: i = 0
for model in model_df['model_dict']:
    i = i + 1
    print('Model', i)
    for param_tensor in model:
        print(param_tensor, "\t", model[param_tensor].size())

    print()

i = 0
for optimizer in model_df['optim_dict']:
    i = i + 1
    print('Model', i)
    for var_name in optimizer:
        if var_name == 'state':
            a = 1
        else:
            print((var_name), "\t", optimizer[var_name])

    print()
```

```
Model 1
fc1.weight      torch.Size([200, 3072])
fc1.bias        torch.Size([200])
fc2.weight      torch.Size([200, 200])
fc2.bias        torch.Size([200])
fc3.weight      torch.Size([10, 200])
fc3.bias        torch.Size([10])
```

```
Model 2
fc1.weight      torch.Size([200, 3072])
fc1.bias        torch.Size([200])
fc2.weight      torch.Size([200, 200])
fc2.bias        torch.Size([200])
fc3.weight      torch.Size([10, 200])
fc3.bias        torch.Size([10])
```

```
Model 3
fc1.weight      torch.Size([200, 3072])
fc1.bias        torch.Size([200])
fc2.weight      torch.Size([200, 200])
fc2.bias        torch.Size([200])
fc3.weight      torch.Size([10, 200])
```



```
fc3.bias          torch.Size([10])
```

#### Model 4

```
fc1.weight        torch.Size([50, 3072])
fc1.bias          torch.Size([50])
fc2.weight        torch.Size([50, 50])
fc2.bias          torch.Size([50])
fc3.weight        torch.Size([10, 50])
fc3.bias          torch.Size([10])
```

#### Model 5

```
fc1.weight        torch.Size([50, 3072])
fc1.bias          torch.Size([50])
fc2.weight        torch.Size([50, 50])
fc2.bias          torch.Size([50])
fc3.weight        torch.Size([10, 50])
fc3.bias          torch.Size([10])
```

#### Model 6

```
fc1.weight        torch.Size([50, 3072])
fc1.bias          torch.Size([50])
fc2.weight        torch.Size([50, 50])
fc2.bias          torch.Size([50])
fc3.weight        torch.Size([10, 50])
fc3.bias          torch.Size([10])
```

#### Model 7

```
fc1.weight        torch.Size([50, 3072])
fc1.bias          torch.Size([50])
fc2.weight        torch.Size([50, 50])
fc2.bias          torch.Size([50])
fc3.weight        torch.Size([10, 50])
fc3.bias          torch.Size([10])
```

#### Model 8

```
fc1.weight        torch.Size([5, 3072])
fc1.bias          torch.Size([5])
fc2.weight        torch.Size([5, 5])
fc2.bias          torch.Size([5])
fc3.weight        torch.Size([10, 5])
fc3.bias          torch.Size([10])
```

#### Model 1

```
param_groups      [{'lr': 0.001, 'momentum': 0.9, 'dampening': 0, 'weight_decay':
0, 'nesterov': False, 'params': [0, 1, 2, 3, 4, 5]}]
```

#### Model 2

```
param_groups      [{'lr': 0.001, 'momentum': 0.9, 'dampening': 0, 'weight_decay':
0.01, 'nesterov': False, 'params': [0, 1, 2, 3, 4, 5]}]
```

#### Model 3

```
param_groups      [{'lr': 0.001, 'momentum': 0.9, 'dampening': 0, 'weight_decay':
0.5, 'nesterov': False, 'params': [0, 1, 2, 3, 4, 5]}]
```

#### Model 4

```
param_groups      [{'lr': 0.001, 'momentum': 0.9, 'dampening': 0, 'weight_decay':
0, 'nesterov': False, 'params': [0, 1, 2, 3, 4, 5]}]
```

#### Model 5

```
param_groups      [{'lr': 0.001, 'momentum': 0.9, 'dampening': 0, 'weight_decay':
0.01, 'nesterov': False, 'params': [0, 1, 2, 3, 4, 5]}]
```

#### Model 6

```
param_groups      [{'lr': 0.001, 'momentum': 0.9, 'dampening': 0, 'weight_decay':
0.5, 'nesterov': False, 'params': [0, 1, 2, 3, 4, 5]}]
```

Model 7

```
param_groups      [{'lr': 0.001, 'momentum': 0.9, 'dampening': 0, 'weight_decay': 0.5, 'nesterov': False, 'params': [0, 1, 2, 3, 4, 5]}]
```

Model 8

```
param_groups      [{'lr': 0.001, 'momentum': 0.9, 'dampening': 0, 'weight_decay': 0, 'nesterov': False, 'params': [0, 1, 2, 3, 4, 5]}]
```

## BONUS. Fun extra tests

```
In [13]: # # check that test data are set up
# dataiter = iter(test_loader)
# images, labels = dataiter.next()

# # print images
# imshow(torchvision.utils.make_grid(images))
# print('GroundTruth: ', ' '.join('%5s' % classes[labels[j]] for j in range(batch_size)))

# # preliminary check
# outputs = net_in(images)
# _, predicted = torch.max(outputs, 1)

# print('Predicted: ', ' '.join('%5s' % classes[predicted[j]]
#                               for j in range(batch_size)))

# # final check
# correct = 0
# total = 0
# with torch.no_grad():
#     for data in test_loader:
#         images, labels = data
#         outputs = net_in(images)
#         _, predicted = torch.max(outputs.data, 1)
#         total += labels.size(0)
#         correct += (predicted == labels).sum().item()

# print('\n Accuracy of the network on the 10000 test images: %d %%' % (
#     100 * correct / total))

# # What are the classes that performed well, and the classes that did not perform well?
# print('')
# class_correct = list(0. for i in range(10))
# class_total = list(0. for i in range(10))
# with torch.no_grad():
#     for data in test_loader:
#         images, labels = data
#         outputs = net_in(images)
#         _, predicted = torch.max(outputs, 1)
#         c = (predicted == labels).squeeze()
#         for i in range(4):
#             label = labels[i]
#             class_correct[label] += c[i].item()
#             class_total[label] += 1

# for i in range(10):
#     print('Accuracy of %5s : %2d %%' % (
#         classes[i], 100 * class_correct[i] / class_total[i]))
```



GroundTruth:    cat   ship   ship plane   frog   frog   car   frog   cat   car plane  
truck

Predicted:    cat truck plane   bird   deer   frog   cat   frog   dog   car plane tr  
uck

Accuracy of the network on the 10000 test images: 53 %

Accuracy of plane : 61 %

Accuracy of car : 61 %

Accuracy of bird : 31 %

Accuracy of cat : 36 %

Accuracy of deer : 44 %

Accuracy of dog : 41 %

Accuracy of frog : 59 %

Accuracy of horse : 66 %

Accuracy of ship : 69 %

Accuracy of truck : 61 %

# HW\_04 Problem #2

## 2 Adaboost [20pts]

Write a class that implements the Adaboost algorithm. Your class should be similar to sklearn's in that it should have a `fit` and `predict` method to train and test the classifier, respectively. You should also use the sampling function from Homework #1 to train the weak learning algorithm, which should be a shallow decision tree. The Adaboost class should be compared to sklearn's implementation on datasets from the course Github page.

This Code was written by Robert 'Quinn' Hull, and borrowed elements from several other resources:

- Understanding AdaBoost - <https://towardsdatascience.com/understanding-adaboost-2f94f22d5bfe>
- SKLearn - <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.AdaBoostClassifier.html>
- Machine Learning Mastery
- SKLearn - [https://scikit-learn.org/stable/auto\\_examples/ensemble/plot\\_adaboost\\_twoclass.html#sphx-glr-auto-examples-ensemble-plot-adaboost-twoclass-py](https://scikit-learn.org/stable/auto_examples/ensemble/plot_adaboost_twoclass.html#sphx-glr-auto-examples-ensemble-plot-adaboost-twoclass-py)

P2

①  $h_t = \text{WeakLearn}(D_t, \{x_i, y_i\}, \mathcal{H})$  tree

②  $\epsilon_t = \sum_{i \in \text{wrong}} D_t(i)$  ↳ 1)  $c = \text{CART}(\text{max\_depth}=4)$

③  $\alpha_t = \frac{1}{2} \log \frac{1 - \epsilon_t}{\epsilon_t}$  2)  $\text{idx} = \text{rsample}(D, \text{len}(D))$  - find indices

④  $D_{t+1}(i) = \frac{D_t(i)}{Z_t} \exp \{ -\alpha_t y_i h_t(x_i) \}$  3)  $X_{t0}, y_t = X[\text{idx}], y[\text{idx}]$  - take subset of x and get their labels

4)  $h_t = c.\text{fit}(X_{t0}, y_t)$  - give hypothesis

How to use the sampling distribution function to get your weak learning classifier

generic classifier

of data 15:04 samples you want to train on

15:14

for  $t=1, \dots, T$  do

creates a hypothesis

①  $h_t = \text{Weak Learn}(S, D_t, \mathcal{H})$

pts that are hard to classify if you get them wrong are going to effect error more

②  $\epsilon_t = \sum_{i=1}^n D_t(i) \mathbb{I}[h(x_i) \neq y_i]$  [Error]

dataset distribution hypothesis

indicator function

$\mathbb{I}[\cdot] = \begin{cases} 1, & \text{true} \\ 0, & \text{false} \end{cases}$

③  $\alpha_t = \frac{1}{2} \log \frac{1-\epsilon_t}{\epsilon_t}$

some sort of update parameter that tells us how good the model is

if error is small, alpha is higher  
if error is large, alpha is small  
nb: alpha can never be negative  
if model is better than a guess

④  $D_{t+1}(i) = \frac{D_t(i)}{Z_t} \exp(-\alpha_t h_t(x_i) y_i)$

norm constant

If wrong, negative  
If right, positive

If we get point right, weight goes down  
If we get point wrong, weight goes up

---

**Algorithm 1** Adaboost (Adaptive Boosting)
 

---

**Input:**  $S := \{x_i, y_i\}_{i=1}^N$ , learning rounds  $T$ , and hypothesis class  $\mathcal{H}$

**Initialize:**  $D_1(i) = 1/n$

- 1: **for**  $t = 1, \dots, T$  **do**
  - 2:  $h_t = \text{WEAKLEARN}(h, S, D_t)$
  - 3:  $\epsilon_t = \sum D_t(i) \mathbb{I}[h(x_i) \neq y_i]$
  - 4:  $\alpha_t = \frac{1}{2} \log \left( \frac{1-\epsilon_t}{\epsilon_t} \right)$
  - 5:  $D_{t+1}(i) = \frac{D_t(i)}{Z_t} \exp(-\alpha_t y_i h_t(x_i))$
  - 6: **end for**
  - 7: **Output:**  $H(x) = \text{sign} \left( \sum_{t=1}^T \alpha_t h_t(x) \right)$
- 

## AdaBoosting

### Pseudocode

- for  $t=1$  to  $T$ :
  - 1) execute a weaklearn hypothesis ( $h_t$ )
    - set generic classifier,  $c = \text{Cart}(\text{max\_depth}=4)$
    - set indices of sample to train,  $\text{idx} = \text{rsample}(D, \text{len}(D))$
    - Take subset of  $x$  and get labels,  $x_t, y_t = X[\text{idx}], y[\text{idx}]$
    - create hypothesis,  $h[t] = c.\text{fit}(x_t, y_t)$
  - 2) calculate the error ( $E_t$ )
  - 3) update change parameter
  - 4) reset weights

### Modules

```
In [1]: import numpy as np
import matplotlib.pyplot as plt

from sklearn.ensemble import AdaBoostClassifier
```

```
from sklearn.tree import DecisionTreeClassifier
from sklearn.datasets import make_gaussian_quantiles
```

## Create Data Sets

```
In [2]: # Construct dataset
X1, y1 = make_gaussian_quantiles(cov=2.,
                                n_samples=200, n_features=2,
                                n_classes=2, random_state=1)
X2, y2 = make_gaussian_quantiles(mean=(3, 3), cov=1.5,
                                n_samples=300, n_features=2,
                                n_classes=2, random_state=1)

X = np.concatenate((X1, X2))
y = np.concatenate((y1, - y2 + 1))
y[(y == 0)] = -1 # reset '0' classes to -1
```

## AdaBoost

### resample

```
In [3]: def random(M_in, p_pdf):
        """
        M_in : the number of samples
        p_pdf : a probability distribution (not cumulative)
        """
        p_cdf = np.cumsum(p_pdf)
        output = np.zeros(M_in)
        for m in range(M_in): # loop through all elements in array
            p_in = np.random.uniform(0,1) # generate a random probability
            p_i = np.where(p_cdf >= p_in)[0][0] # extract first index of condition
            output[m] = p_i # randomly generate a number in probability interview
        return output
```

## WeakLearn

```
In [4]: def weaklearn(H, X, y, D):
        """
        H : Hypothesis Class
        D : Weight Distribution
        X, y : Dataset
        """
        # set generic classifier
        c = H
        # get indices of samples to train
        idx = random(M_in=len(D), p_pdf=D).astype(np.int16)
        # take subset of x and get labels
        xt, yt = X[idx], y[idx]
        # create hypothesis
        ht = c.fit(xt, yt)
        return ht
```

## Error

```
In [5]: def Error(D,ht,X,y):
        """
        D : distribution of weights
```

```

ht : weak learning classifier
X : array of X values (data set)
y : array of y values (data set)
...

err_t = 0
for i in range(len(D)):
    # print((ht.predict([X[i]]) != [y[i]]))
    err_i = D[i]*(ht.predict([X[i]]) != [y[i]])
    err_t = err_t + err_i

return err_t

```

## Adaboost Code - homegrown training

```

In [6]: # Globals
T = 50 # number of rounds
depth = 50 # number of estimators
H = DecisionTreeClassifier(max_depth=depth) # h is a classifier int he hypothesis
n = len(y) # number of samples in data set
N_arr = np.arange(0,n,1) # an array of the different indexes

```

```

In [7]: err_list = [] # for keeping track of errors
Dt = np.ones(n)*(1/n) # initializing weights, as equal to 1 / number of samples.
# print(Dt.sum())
for t in range(T):
    # generate hypothesis
    ht = weaklearn(H=H,X=X,y=y,D=Dt)
    # print(ht)

    # calculate error
    err_t = Errorrt(D=Dt,ht=ht,X=X,y=y)
    err_list.append(err_t)
    # print('error is', err_t)

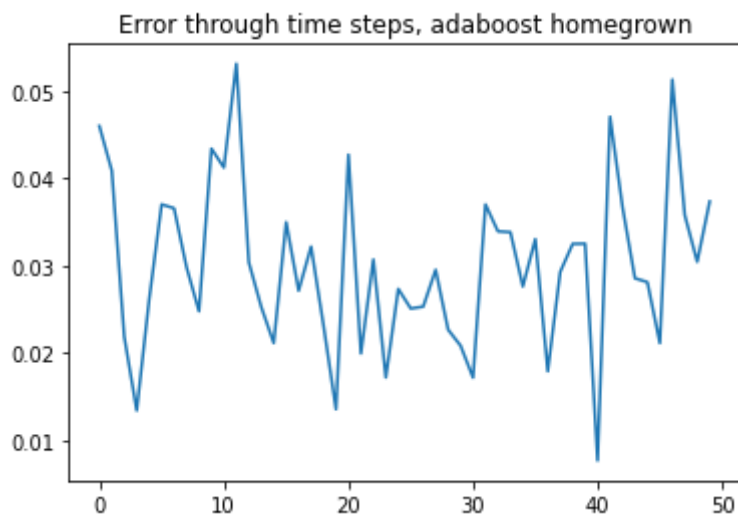
    # calculate update parameter - low error, high alpha
    alpha_t = 0.5*np.log((1-err_t)/err_t)
    # print('alpha is', alpha_t)

    # calculate norm constant - a function of the error, according to lecture note
    Zt = 2*np.sqrt(err_t*(1-err_t))
    # print(Zt)

    # update weights
    Dt = (Dt / Zt) * np.exp(-alpha_t*ht.predict(X)*y)
    # print('Weights summed to', Dt.sum())

plt.plot(np.arange(0,T,1),err_list)
plt.title('Error through time steps, adaboost homegrown')
plt.show()

```



## TESTING: Adaboost Code - homegrown

```
In [9]: y[(y == 0)] = -1

ht.fit(X,y)

plot_step = 0.02

plt.figure(figsize=(10, 5))

# Plot the decision boundaries
plt.subplot(121)
x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, plot_step),
                     np.arange(y_min, y_max, plot_step))


Z = ht.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)
cs = plt.contourf(xx, yy, Z, cmap=plt.cm.Paired)
plt.axis("tight")

# Plot the training points
for i in range(-1,2,2):
    # print(i)
    idx = np.where(y == i)
    # print(idx)
    plt.scatter(X[idx, 0], X[idx, 1],
                cmap=plt.cm.Paired,
                s=20, edgecolor='k',
                label="Class %s" % str(i))
plt.xlim(x_min, x_max)
plt.ylim(y_min, y_max)
plt.legend(loc='upper right')
plt.xlabel('x')
plt.ylabel('y')
plt.title('Decision Boundary for homegrown implementation')
```

Out[9]: Text(0.5, 1.0, 'Decision Boundary for homegrown implementation')



Decision Boundary for homegrown implementation



## TESTING: Adaboost - From SKLearn

```
In [10]: print(__doc__)

# Author: Noel Dawe <noel.dawe@gmail.com>
#
# License: BSD 3 clause

y[(y == -1)] = 0

# Create and fit an AdaBoosted decision tree
bdt = AdaBoostClassifier(DecisionTreeClassifier(max_depth=1),
                        algorithm="SAMME",
                        n_estimators=depth)

bdt.fit(X, y)

plot_colors = "br"
plot_step = 0.02
class_names = "AB"

plt.figure(figsize=(10, 5))

# Plot the decision boundaries
plt.subplot(121)
x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, plot_step),
                    np.arange(y_min, y_max, plot_step))


Z = bdt.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)
cs = plt.contourf(xx, yy, Z, cmap=plt.cm.Paired)
plt.axis("tight")

# Plot the training points
for i, n, c in zip(range(2), class_names, plot_colors):
    idx = np.where(y == i)
```

```
plt.scatter(X[idx, 0], X[idx, 1],
            c=c, cmap=plt.cm.Paired,
            s=20, edgecolor='k',
            label="Class %s" % n)
plt.xlim(x_min, x_max)
plt.ylim(y_min, y_max)
plt.legend(loc='upper right')
plt.xlabel('x')
plt.ylabel('y')
plt.title('Decision Boundary For Sklearn implementation')
```

Automatically created module for IPython interactive environment

Out[10]: Text(0.5, 1.0, 'Decision Boundary For Sklearn implementation')



In [ ]:

# HW\_04 Problem #3

## 3 Recurrent Neural Networks for Language Modeling [20pts]

Read “LSTM: A Search Space Odyssey” (<https://arxiv.org/abs/1503.04069>). One application of an RNN is the ability model language, which is what your phone does when it is predicting the top three words when you’re texting. In this problem, you will need to build a language model.

You are encouraged to start out with the code [here](#). While this code will implement a language model, you are required to modify the code to attempt to beat the baseline for the experiments they have implemented. For example, one modification would be to train multiple language models and average, or weight, their outputs to generate language. Write a couple of paragraphs about what you did and if the results improve the model over the baseline on Github.

This Code was written by Robert 'Quinn' Hull, and borrowed elements from several other resources:

- The bulk of this script is the TensorFlow article about text generation: [https://www.tensorflow.org/tutorials/text/text\\_generation](https://www.tensorflow.org/tutorials/text/text_generation)
- Much of the text was removed to make this shorter and more intelligible
- The text describing my work is available at the end of the script.

```
In [212...  ##@title Licensed under the Apache License, Version 2.0 (the "License");
## you may not use this file except in compliance with the License.
## You may obtain a copy of the License at
##
## https://www.apache.org/licenses/LICENSE-2.0
##
## Unless required by applicable law or agreed to in writing, software
## distributed under the License is distributed on an "AS IS" BASIS,
## WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
## See the License for the specific language governing permissions and
## limitations under the License.
```

## Setup

### Import TensorFlow and other libraries

```
In [265...  import tensorflow as tf
from tensorflow.keras.layers.experimental import preprocessing

import numpy as np
import os
import time
```

### Download the Shakespeare dataset

```
In [266...  path_to_file = tf.keras.utils.get_file('shakespeare.txt', 'https://storage.googl
```

### Read the data

First, look in the text:

```
In [267... # Read, then decode for py2 compat.
text = open(path_to_file, 'rb').read().decode(encoding='utf-8')
# # subset training text (temporary)
# text = text[:11153]
# length of text is the number of characters in it
print('Length of text: {} characters'.format(len(text)))
# The unique characters in the file
vocab = sorted(set(text))
print('{} unique characters'.format(len(vocab)))
```

Length of text: 1115394 characters  
65 unique characters

## Process the text

### Vectorize the text

```
In [268... example_texts = ['abcdefg', 'xyz']

chars = tf.strings.unicode_split(example_texts, input_encoding='UTF-8')
print('characters in ', chars)

ids_from_chars = preprocessing.StringLookup(
    vocabulary=list(vocab))

ids = ids_from_chars(chars)
ids

print('IDs ', ids)

chars_from_ids = tf.keras.layers.experimental.preprocessing.StringLookup(
    vocabulary=ids_from_chars.get_vocabulary(), invert=True)

chars = chars_from_ids(ids)
chars

print('characters out ', chars)

# You can `tf.strings.reduce_join` to join the characters back into strings.
def text_from_ids(ids):
    return tf.strings.reduce_join(chars_from_ids(ids), axis=-1)

characters in <tf.RaggedTensor [[b'a', b'b', b'c', b'd', b'e', b'f', b'g'],
[b'x', b'y', b'z']]>
IDs <tf.RaggedTensor [[41, 42, 43, 44, 45, 46, 47], [64, 65, 66]]>
characters out <tf.RaggedTensor [[b'a', b'b', b'c', b'd', b'e', b'f', b'g'],
[b'x', b'y', b'z']]>
```

### Create training examples, targets, batches

```
In [269... # Create examples and targets
all_ids = ids_from_chars(tf.strings.unicode_split(text, 'UTF-8'))
ids_dataset = tf.data.Dataset.from_tensor_slices(all_ids)

# set constants
seq_length = 100
examples_per_epoch = len(text)//(seq_length+1)
```

```

# set batches
sequences = ids_dataset.batch(seq_length+1, drop_remainder=True)

# split dataset
def split_input_target(sequence):
    input_text = sequence[:-1]
    target_text = sequence[1:]
    return input_text, target_text

dataset = sequences.map(split_input_target)

# create training batches
# Batch size
BATCH_SIZE = 64

# Buffer size to shuffle the dataset
# (TF data is designed to work with possibly infinite sequences,
# so it doesn't attempt to shuffle the entire sequence in memory. Instead,
# it maintains a buffer in which it shuffles elements).
BUFFER_SIZE = 10000

dataset = (
    dataset
    .shuffle(BUFFER_SIZE)
    .batch(BATCH_SIZE, drop_remainder=True)
    .prefetch(tf.data.experimental.AUTOTUNE))

dataset

```

Out[269... <PrefetchDataset shapes: ((64, 100), (64, 100)), types: (tf.int64, tf.int64)>

## Build The Model

```

In [285... # models
class MyModel(tf.keras.Model):
    def __init__(self, vocab_size, embedding_dim, rnn_units):
        super().__init__(self)
        self.embedding = tf.keras.layers.Embedding(vocab_size, embedding_dim)
        self.gru = tf.keras.layers.GRU(rnn_units,
                                         return_sequences=True,
                                         return_state=True)
        self.dense = tf.keras.layers.Dense(vocab_size)

    def call(self, inputs, states=None, return_state=False, training=False):
        x = inputs
        x = self.embedding(x, training=training)
        if states is None:
            states = self.gru.get_initial_state(x)
        x, states = self.gru(x, initial_state=states, training=training)
        x = self.dense(x, training=training)

        if return_state:
            return x, states
        else:
            return x

# new models, an LSTM
class NewModel1(tf.keras.Model):

```

```

def __init__(self, vocab_size, embedding_dim, rnn_units):
    super().__init__(self)
    self.embedding = tf.keras.layers.Embedding(vocab_size, embedding_dim)
    self.lstm = tf.keras.layers.LSTM(rnn_units,
                                     return_sequences=True,
                                     return_state=False) # change from GRU to LSTM
    self.dense = tf.keras.layers.Dense(vocab_size)

def call(self, inputs, states=None, return_state=False, training=False):
    x = inputs
    x = self.embedding(x, training=training)
    # print(states)
    # print(training)
    if states is None:
        states = self.lstm.get_initial_state(x)
    # print(states)
    x = self.lstm(x) # , initial_state=states) # , training=training)
    x = self.dense(x, training=training)

    if return_state:
        return x, states
    else:
        return x

# new models2, an LSTM
class NewModel2(tf.keras.Model):
    def __init__(self, vocab_size, embedding_dim, rnn_units):
        super().__init__(self)
        self.embedding = tf.keras.layers.Embedding(vocab_size, embedding_dim)
        self.lstm = tf.keras.layers.LSTM(rnn_units,
                                         return_sequences=True,
                                         return_state=True) # change from GRU to LSTM
        self.dense = tf.keras.layers.Dense(vocab_size)

    def call(self, inputs, states=None, return_state=False, training=False):
        x = inputs
        x = self.embedding(x, training=training)
        # print(states)
        # print(training)
        if states is None:
            states = self.lstm.get_initial_state(x)
        # print(states)
        x = self.lstm(x, initial_state=states, training=training)
        x = self.dense(x, training=training)

        if return_state:
            return x, states
        else:
            return x

```

In [286...

```

# model chars
# Length of the vocabulary in chars
vocab_size = len(vocab) # QH NOTE: this = 65.

# The embedding dimension
embedding_dim = 256

# Number of RNN units
rnn_units = int(1024) # 1024

```

```
# set model
model = NewModel2(
    # Be sure the vocabulary size matches the `StringLookup` layers.
    vocab_size=len(ids_from_chars.get_vocabulary()),
    embedding_dim=embedding_dim,
    rnn_units=rnn_units)

# optimizer and loss function
loss = tf.losses.SparseCategoricalCrossentropy(from_logits=True)
model.compile(optimizer='adam', loss=loss)

# set number of epochs
EPOCHS = 10
```

## Training

### Easy Train

```
In [287... # Directory where the checkpoints will be saved
checkpoint_dir = './training_checkpoints'
# Name of the checkpoint files
checkpoint_prefix = os.path.join(checkpoint_dir, "ckpt_{epoch}")

checkpoint_callback = tf.keras.callbacks.ModelCheckpoint(
    filepath=checkpoint_prefix,
    save_weights_only=True)

# train model (naive)
history = model.fit(dataset, epochs=EPOCHS, callbacks=[checkpoint_callback])
```

Epoch 1/10

```
-----
ValueError                                Traceback (most recent call last)
<ipython-input-287-28b3bf9f2177> in <module>()
      9
     10 # train model (naive)
--> 11 history = model.fit(dataset, epochs=EPOCHS, callbacks=[checkpoint_callba
ck])

/usr/local/lib/python3.7/dist-packages/tensorflow/python/keras/engine/training.p
y in fit(self, x, y, batch_size, epochs, verbose, callbacks, validation_split, v
alidation_data, shuffle, class_weight, sample_weight, initial_epoch, steps_per_e
poch, validation_steps, validation_batch_size, validation_freq, max_queue_size,
workers, use_multiprocessing)
    1098         _r=1):
    1099             callbacks.on_train_batch_begin(step)
--> 1100             tmp_logs = self.train_function(iterator)
    1101             if data_handler.should_sync:
    1102                 context.async_wait()

/usr/local/lib/python3.7/dist-packages/tensorflow/python/eager/def_function.py i
n __call__(self, *args, **kwargs)
    826         tracing_count = self.experimental_get_tracing_count()
    827         with trace.Trace(self._name) as tm:
--> 828             result = self._call(*args, **kwargs)
    829             compiler = "xla" if self._experimental_compile else "nonXla"
    830             new_tracing_count = self.experimental_get_tracing_count()

/usr/local/lib/python3.7/dist-packages/tensorflow/python/eager/def_function.py i
n __call__(self, *args, **kwargs)
    869         # This is the first call of __call__, so we have to initialize.
```

```

870         initializers = []
--> 871         self._initialize(args, kwds, add_initializers_to=initializers)
872     finally:
873         # At this point we know that the initialization is complete (or le
ss

/usr/local/lib/python3.7/dist-packages/tensorflow/python/eager/def_function.py i
n _initialize(self, args, kwds, add_initializers_to)
874         self._concrete_stateful_fn = (
875             self._stateful_fn._get_concrete_function_internal_garbage_collec
ted( # pylint: disable=protected-access
--> 876                 *args, **kwds))
877
878     def invalid_creator_scope(*unused_args, **unused_kwds):

/usr/local/lib/python3.7/dist-packages/tensorflow/python/eager/function.py in _g
et_concrete_function_internal_garbage_collected(self, *args, **kwargs)
2967         args, kwargs = None, None
2968         with self._lock:
-> 2969             graph_function, _ = self._maybe_define_function(args, kwargs)
2970         return graph_function
2971

/usr/local/lib/python3.7/dist-packages/tensorflow/python/eager/function.py in _m
aybe_define_function(self, args, kwargs)
3359
3360         self._function_cache.missed.add(call_context_key)
-> 3361         graph_function = self._create_graph_function(args, kwargs)
3362         self._function_cache.primary[cache_key] = graph_function
3363

/usr/local/lib/python3.7/dist-packages/tensorflow/python/eager/function.py in _c
reate_graph_function(self, args, kwargs, override_flat_arg_shapes)
3204         arg_names=arg_names,
3205         override_flat_arg_shapes=override_flat_arg_shapes,
-> 3206         capture_by_value=self._capture_by_value),
3207         self._function_attributes,
3208         function_spec=self.function_spec,

/usr/local/lib/python3.7/dist-packages/tensorflow/python/framework/func_graph.py
in func_graph_from_py_func(name, python_func, args, kwargs, signature, func_grap
h, autograph, autograph_options, add_control_dependencies, arg_names, op_return_
value, collections, capture_by_value, override_flat_arg_shapes)
988         _, original_func = tf_decorator.unwrap(python_func)
989
--> 990         func_outputs = python_func(*func_args, **func_kwargs)
991
992         # invariant: `func_outputs` contains only Tensors, CompositeTensor
s,

/usr/local/lib/python3.7/dist-packages/tensorflow/python/eager/def_function.py i
n wrapped_fn(*args, **kwds)
632         xla_context.Exit()
633     else:
--> 634         out = weak_wrapped_fn().__wrapped__(*args, **kwds)
635         return out
636

/usr/local/lib/python3.7/dist-packages/tensorflow/python/framework/func_graph.py
in wrapper(*args, **kwargs)
975         except Exception as e: # pylint:disable=broad-except
976             if hasattr(e, "ag_error_metadata"):
--> 977                 raise e.ag_error_metadata.to_exception(e)
978             else:
979                 raise

```



**ValueError:** in user code:

```

/usr/local/lib/python3.7/dist-packages/tensorflow/python/keras/engine/training.py:805 train_function *
    return step_function(self, iterator)
<ipython-input-279-ea55be33f8d6>:69 call *
    x = self.dense(x, training=training)
/usr/local/lib/python3.7/dist-packages/tensorflow/python/keras/engine/base_layer.py:998 __call__ **
    input_spec.assert_input_compatibility(self.input_spec, inputs, self.name)
/usr/local/lib/python3.7/dist-packages/tensorflow/python/keras/engine/input_spec.py:207 assert_input_compatibility
    'input tensors. Inputs received: ' + str(inputs))

```

ValueError: Layer dense\_28 expects 1 input(s), but it received 3 input tensors. Inputs received: [<tf.Tensor 'new\_model2\_1/lstm\_15/PartitionedCall:1' shape=(64, 100, 1024) dtype=float32>, <tf.Tensor 'new\_model2\_1/lstm\_15/PartitionedCall:2' shape=(64, 1024) dtype=float32>, <tf.Tensor 'new\_model2\_1/lstm\_15/PartitionedCall:3' shape=(64, 1024) dtype=float32>]

## Generate text

The following makes a single step prediction:

In [273...

```

class OneStep(tf.keras.Model):
    def __init__(self, model, chars_from_ids, ids_from_chars, temperature=1.0):
        super().__init__()
        self.temperature=temperature
        self.model = model
        self.chars_from_ids = chars_from_ids
        self.ids_from_chars = ids_from_chars

        # Create a mask to prevent "" or "[UNK]" from being generated.
        skip_ids = self.ids_from_chars(['', '[UNK]'][:, None])
        sparse_mask = tf.SparseTensor(
            # Put a -inf at each bad index.
            values=[-float('inf')]*len(skip_ids),
            indices = skip_ids,
            # Match the shape to the vocabulary
            dense_shape=[len(ids_from_chars.get_vocabulary())])
        self.prediction_mask = tf.sparse.to_dense(sparse_mask)

    @tf.function
    def generate_one_step(self, inputs, states=None):
        # Convert strings to token IDs.
        input_chars = tf.strings.unicode_split(inputs, 'UTF-8')
        input_ids = self.ids_from_chars(input_chars).to_tensor()

        # Run the model.
        # predicted_logits.shape is [batch, char, next_char_logits]
        predicted_logits, states = self.model(inputs=input_ids, states=states,
                                              return_state=True)

        # Only use the last prediction.
        predicted_logits = predicted_logits[:, -1, :]
        predicted_logits = predicted_logits/self.temperature
        # Apply the prediction mask: prevent "" or "[UNK]" from being generated.
        predicted_logits = predicted_logits + self.prediction_mask

        # Sample the output logits to generate token IDs.

```

```

predicted_ids = tf.random.categorical(predicted_logits, num_samples=1)
predicted_ids = tf.squeeze(predicted_ids, axis=-1)

# Convert from token ids to characters
predicted_chars = self.chars_from_ids(predicted_ids)

# Return the characters and model state.
return predicted_chars, states

```

```

In [274... temp = 1.0
one_step_model = OneStep(model, chars_from_ids, ids_from_chars, temperature=temp

```

```

In [275... start = time.time()
states = None
next_char = tf.constant(['ROMEO:'])
result = [next_char]

pred_char_len = 50

for n in range(pred_char_len):
    next_char, states = one_step_model.generate_one_step(next_char, states=states)
    result.append(next_char)

result = tf.strings.join(result)
end = time.time()

print(result[0].numpy().decode('utf-8'), '\n\n' + '_'*80)

print(f"\nRun time: {end - start}")

```

ROMEO:  
 Ag:  
 Ange notifaloulyo ced bleathird thostll?  
 Thit

---

Run time: 0.8481040000915527

## Evaluate

### Bleu

- From: <https://towardsdatascience.com/how-to-evaluate-text-generation-models-metrics-for-automatic-evaluation-of-nlp-models-e1c251b04ec1>
- BLEU is a precision focused metric that calculates n-gram overlap of the reference and generated texts.

```

In [276... # prep
from nltk.translate.bleu_score import SmoothingFunction, corpus_bleu, sentence_bleu

def bleu(ref, gen):
    """
    calculate pair wise bleu score. uses nltk implementation
    Args:
        For word comparison
        references : a list of reference sentences
        candidates : a list of candidate(generated) sentences
    """

```

```

    For character comparison
        references : a list of reference sentences
        candidates : a list of candidate(generated) sentences
Returns:
    bleu score(float)
'''
ref_bleu = []
gen_bleu = []
print('test1')
for l in gen:
    print('test2')
    gen_bleu.append(l.split())
print(gen_bleu)
for i,l in enumerate(ref):
    print('test3')
    ref_bleu.append([l.split()])
print(ref_bleu)
cc = SmoothingFunction()
print(cc)
score_bleu = corpus_bleu(ref_bleu, gen_bleu)
return score_bleu

```

```
In [277...] bleu_score = 'NA' # bleu(['hillo'], ['hello'])
```

## Final Comments and Text

- In general, the script experiments with:
  - hyperparameters:
    - embedding\_dim, rnn\_units, temperature, epochs
  - Model Structure
    - LSTM v RNN
- Lingering questions I have:
  - I was mostly unsuccessful at using an LSTM for this. It's definitely in how I am setting up this model.
  - I never figured out how to properly assess this model. I started to experiment with the BLEU algorithm, but I wasn't sure this is built for character generation in this context. I.E., I wasn't sure how to apply this to our use-case where the text generated was mostly 'random'

```
In [278...] cats = ['bleu_score', 'Training loss (final)', 'Training sequence length',
                  'Training Buffer Size', 'Training Epochs', 'Training Batch Size',
                  'RNN model: RNN units', 'RNN model: Embedding Dim',
                  'RNN model: Name of loss function', 'RNN model: Summary', 'RNN model: hi
                  'Prediction: character length', 'Prediction: temp constant', 'Prediction
                  'Prediction: model object']

print(cats, '\n')

# save1 = [bleu_score, history.history['loss'][-1], seq_length,
#          BUFFER_SIZE, EPOCHS, BATCH_SIZE,
#          rnn_units, embedding_dim,
#          loss.name, model, history,
#          pred_char_len, temp, result,
#          one_step_model]
```

```

print('Model 1 is the `custom` model using the customized model object in the or
      'There is no calculation of bleu score for this one, because I havent figu
      'how to best evaluate the model. \n')
print(save1, '\n\n')

# save2 = [bleu_score, history.history['loss'][-1], seq_length,
#          BUFFER_SIZE, EPOCHS, BATCH_SIZE,
#          rnn_units, embedding_dim,
#          loss.name, model, history,
#          pred_char_len, temp, result,
#          one_step_model]
print('Model 2 is our baseline model using the class MyModel with architecture t
      'From the perspective of training loss, it performs nearly as well as the p
      np.round(save2[1],2), 'versus', np.round(save1[1],2), '\n')
print(save2, '\n\n')

# save3 = [bleu_score, history.history['loss'][-1], seq_length,
#          BUFFER_SIZE, EPOCHS, BATCH_SIZE,
#          rnn_units, embedding_dim,
#          loss.name, model, history,
#          pred_char_len, temp, result,
#          one_step_model]
print('Model 3 is our baseline model using the class MyModel with architecture t
      'to highlight the impact of epochs on training performance, we have reduce
      'From the perspective of training loss, it performs worse than before',np.r
      'The predicted text is less coherent, too:\n', 'Model 3 :', save3[13].numpy
print(save3, '\n\n')

# save4 = [bleu_score, history.history['loss'][-1], seq_length,
#          BUFFER_SIZE, EPOCHS, BATCH_SIZE,
#          rnn_units, embedding_dim,
#          loss.name, model, history,
#          pred_char_len, temp, result,
#          one_step_model]
print('Model 4 is our baseline model using the class MyModel with architecture t
      'to highlight the impact of epochs on training performance, we have increa
      'From the perspective of training loss, it performs better than before',np.
      'The predicted text is way more coherent, almost Shakespearean:\n', 'Model
      'This performance bump does come at the expense of time, though (20s / epo
print(save4, '\n\n')

# save5 = [bleu_score, history.history['loss'][-1], seq_length,
#          BUFFER_SIZE, EPOCHS, BATCH_SIZE,
#          rnn_units, embedding_dim,
#          loss.name, model, history,
#          pred_char_len, temp, result,
#          one_step_model]
print('Model 5 is our baseline model using the class MyModel with architecture t
      'to explore ways to speed up training and preserve text coherence, I reduc
      'by two orders of magnitude: 1115394 characters to 11153 characters and ke
      'From the perspective of training loss, it performs far worse than before',
      'The predicted text is nonsense :\n', 'Model 6 :', save5[13].numpy()[0], '
      'This might be because of undertraining, or an issue in how weve indexed t
      'This saves training speed substantially! - 1 s / epoch \n')
print(save5, '\n\n')

# save6 = [bleu_score, history.history['loss'][-1], seq_length,
#          BUFFER_SIZE, EPOCHS, BATCH_SIZE,
#          rnn_units, embedding_dim,
#          loss.name, model, history,
#          pred_char_len, temp, result,

```

```

# one_step_model]
print('Model 6 is our baseline model using the class MyModel, returning to epoch
      'Here we start to vary the architecture of the model, so rnn_units =', save
      'From the perspective of training loss, it performs better than baseline (M
      'The predicted text is probably more coherent :\n', 'Model 6 :', save6[13]
      'Worth noting that the run-time nearly tripled, from 20 s / epoch to 55 s
print(save6, '\n\n')

# save7 = [bleu_score, history.history['loss'][-1], seq_length,
#          BUFFER_SIZE, EPOCHS, BATCH_SIZE,
#          rnn_units, embedding_dim,
#          loss.name, model, history,
#          pred_char_len, temp, result,
#          one_step_model]
print('Model 7 is our baseline model using the class MyModel \n',
      'Where we decrease the rnn_units =', save7[6], 'down from', save6[6], '\n'
      'From the perspective of training loss, it performs worse than baseline (Mo
      'The predicted text is arguably no more or less coherent :\n', 'Model 7 :
      'Worth noting that the run-time didnt change very much \n ')
print(save7, '\n\n')

# save8 = [bleu_score, history.history['loss'][-1], seq_length,
#          BUFFER_SIZE, EPOCHS, BATCH_SIZE,
#          rnn_units, embedding_dim,
#          loss.name, model, history,
#          pred_char_len, temp, result,
#          one_step_model]
print('Model 8 is our first LSTM model using the class NewModel1 \n',
      'Where we return the rnn_units =', save8[6], '\n',
      'From the perspective of training loss, it performs _____ than baseline (Mo
      'The predicted text is totally incoherent :\n', 'Model 8 :', save8[13].num
      'I have definitely made an error in how I set up this LSTM network \n ')
print(save8, '\n\n')

save9 = [bleu_score, history.history['loss'][-1], seq_length,
         BUFFER_SIZE, EPOCHS, BATCH_SIZE,
         rnn_units, embedding_dim,
         loss.name, model, history,
         pred_char_len, temp, result,
         one_step_model]
print('Model 9 is our 2nd LSTM model using the class NewModel2 \n',
      'Ive tried to mess around with the model so that it passes the state\n',
      'From the perspective of training loss, it performs _____ than baseline (Mo
      'The predicted text is totally incoherent :\n', 'Model 9 :', save9[13].num
      'I have definitely made an error in how I set up this LSTM network \n ')
print(save9, '\n\n')

```

```

['bleu_score', 'Training loss (final)', 'Training sequence length', 'Training Bu
ffer Size', 'Training Epochs', 'Training Batch Size', 'RNN model: RNN units', 'R
NN model: Embedding Dim', 'RNN model: Name of loss function', 'RNN model: Summar
y', 'RNN model: history object', 'Prediction: character length', 'Prediction: te
mp constant', 'Prediction: result', 'Prediction: model object']

```

Model 1 is the `custom` model using the customized model object in the original code

There is no calculation of bleu score for this one, because I havent figured out at this point

how to best evaluate the model.

```

['NA', 1.1910383701324463, 100, 10000, 10, 64, 1024, 256, 'sparse_categorical_cr
ossentropy', <__main__.CustomTraining object at 0x7efee66c6cd0>, <tensorflow.pyt
hon.keras.callbacks.History object at 0x7eff4b279410>, 50, 1.0, <tf.Tensor: shap

```

```
e=(1,), dtype=string, numpy=
array([b'ROMEO:\nO, thy sun doth let\nRend her he should he knife.\n'],
      dtype=object)>, <__main__.OneStep object at 0x7efea8986a50>]
```

Model 2 is our baseline model using the class MyModel with architecture the same as in the original code

From the perspective of training loss, it performs nearly as well as the previous custom model

1.2 versus 1.19

```
['NA', 1.2036826610565186, 100, 10000, 10, 64, 1024, 256, 'sparse_categorical_crossentropy', <__main__.MyModel object at 0x7efee66da950>, <tensorflow.python.keras.callbacks.History object at 0x7efee63f06d0>, 50, 1.0, <tf.Tensor: shape=(1,), dtype=string, numpy=
array([b'ROMEO:\nUntim life is he not to guess.\n\nMARCIUS:\nIs there'],
      dtype=object)>, <__main__.OneStep object at 0x7efea7e72810>]
```

Model 3 is our baseline model using the class MyModel with architecture the same as in the original code

to highlight the impact of epochs on training performance, we have reduced it from 10 to 2

From the perspective of training loss, it performs worse than before 2.0 versus 1.2

The predicted text is less coherent, too:

Model 3 : b'ROMEO:\nThe wontely he soum, as\nI sape!\nAnd swant wean I '

Model 2 : b'ROMEO:\nUntim life is he not to guess.\n\nMARCIUS:\nIs there'

```
['NA', 2.0047507286071777, 100, 10000, 2, 64, 1024, 256, 'sparse_categorical_crossentropy', <__main__.MyModel object at 0x7efea7c7bf10>, <tensorflow.python.keras.callbacks.History object at 0x7efea5fa34d0>, 50, 1.0, <tf.Tensor: shape=(1,), dtype=string, numpy=
array([b'ROMEO:\nThe wontely he soum, as\nI sape!\nAnd swant wean I '],
      dtype=object)>, <__main__.OneStep object at 0x7efea7986b10>]
```

Model 4 is our baseline model using the class MyModel with architecture the same as in the original code

to highlight the impact of epochs on training performance, we have increased it from 2 to 30

From the perspective of training loss, it performs better than before 0.48 versus 2.0

The predicted text is way more coherent, almost Shakespearean:

Model 4 : b'ROMEO:\nThou art well where you will hear none.\n\nGREMIO:\n'

Model 3 : b'ROMEO:\nThe wontely he soum, as\nI sape!\nAnd swant wean I '

This performance bump does come at the expense of time, though (20s / epoch \* 30 epochs = 10 minutes,

```
['NA', 0.48344069719314575, 100, 10000, 30, 64, 1024, 256, 'sparse_categorical_crossentropy', <__main__.MyModel object at 0x7efea7090f50>, <tensorflow.python.keras.callbacks.History object at 0x7efea5fa3250>, 50, 1.0, <tf.Tensor: shape=(1,), dtype=string, numpy=
array([b'ROMEO:\nThou art well where you will hear none.\n\nGREMIO:\n'],
      dtype=object)>, <__main__.OneStep object at 0x7efea64a4e90>]
```

Model 5 is our baseline model using the class MyModel with architecture the same as in the original code

to explore ways to speed up training and preserve text coherence, I reduce the size of the input text

by two orders of magnitude: 1115394 characters to 11153 characters and keep epochs = 30

From the perspective of training loss, it performs far worse than before 3.12 versus 0.48

The predicted text is nonsense :

```
Model 6 : b"ROMEO:hEOl',hbta:zNji t atat sAk :nalrt rey e urn m la t"
Model 5 : b'ROMEO:\nThou art well where you will hear none.\n\nGREMIO:\n'
This might be because of undertraining, or an issue in how weve indexed the vocabulary
This saves training speed substantially! - 1 s / epoch
```

```
['NA', 3.115530014038086, 100, 10000, 30, 64, 1024, 256, 'sparse_categorical_crossentropy', <__main__.MyModel object at 0x7efea63b7c10>, <tensorflow.python.keras.callbacks.History object at 0x7efea6362390>, 50, 1.0, <tf.Tensor: shape=(1,), dtype=string, numpy=
array([b"ROMEO:hEOl',hbta:zNji t atat sAk :nalrt rey e urn m la t"]],
      dtype=object)>, <__main__.OneStep object at 0x7efea60b6850>]
```

Model 6 is our baseline model using the class MyModel, returning to epochs = 10 and characters = 1115394

Here we start to vary the architecture of the model, so rnn\_units = 2048 up from 1024

From the perspective of training loss, it performs better than baseline (Model 2) 0.97 versus 1.2

The predicted text is probably more coherent :

```
Model 6 : b'ROMEO:\nThou art not half; he hath not shown but health\na'
```

```
Model 2 : b'ROMEO:\nUntim life is he not to guess.\n\nMARCIVS:\nIs there'
```

Worth noting that the run-time nearly tripled, from 20 s / epoch to 55 s / epoch

```
['NA', 0.9703238010406494, 100, 10000, 10, 64, 2048, 256, 'sparse_categorical_crossentropy', <__main__.MyModel object at 0x7efea4bf0c50>, <tensorflow.python.keras.callbacks.History object at 0x7efea4c0a310>, 50, 1.0, <tf.Tensor: shape=(1,), dtype=string, numpy=
array([b'ROMEO:\nThou art not half; he hath not shown but health\na'],
      dtype=object)>, <__main__.OneStep object at 0x7efea6d12150>]
```

Model 7 is our baseline model using the class MyModel

Where we decrease the rnn\_units = 256 down from 2048

From the perspective of training loss, it performs worse than baseline (Model 2) 1.43 versus 1.2

The predicted text is arguably no more or less coherent :

```
Model 7 : b"ROMEO:\nO! therefaleness of 'I thremfort\nWhere is within "
```

```
Model 2 : b'ROMEO:\nUntim life is he not to guess.\n\nMARCIVS:\nIs there'
```

Worth noting that the run-time didnt change very much

```
['NA', 1.4263566732406616, 100, 10000, 10, 64, 256, 256, 'sparse_categorical_crossentropy', <__main__.MyModel object at 0x7efe42a21290>, <tensorflow.python.keras.callbacks.History object at 0x7efe42bf8c90>, 50, 1.0, <tf.Tensor: shape=(1,), dtype=string, numpy=
array([b"ROMEO:\nO! therefaleness of 'I thremfort\nWhere is within "],
      dtype=object)>, <__main__.OneStep object at 0x7efef3be2b10>]
```

Model 8 is our first LSTM model using the class NewModel1

Where we decrease the rnn\_units = 1024

From the perspective of training loss, it performs \_\_\_\_\_ than baseline (Model 2) 1.31 versus 1.2

The predicted text is totally incoherent :

```
Model 8 : b'ROMEO:\nAg:\nAnge notifaloullyo ced bleathird thostll?\nThit'
```

I have definitely made an error in how I set up this recurrent network

```
['NA', 1.313532829284668, 100, 10000, 10, 64, 1024, 256, 'sparse_categorical_crossentropy', <__main__.NewModel1 object at 0x7efef3f3a090>, <tensorflow.python.keras.callbacks.History object at 0x7efef3ffc090>, 50, 1.0, <tf.Tensor: shape=(1,), dtype=string, numpy=
array([b'ROMEO:\nAg:\nAnge notifaloullyo ced bleathird thostll?\nThit'],
```

```
dtype=object)>, <__main__.OneStep object at 0x7efe42a25810>]
```

In [278...